

Sanity-checking LLM ability scores against agent scores with a fixed scaffold

